

Real-time and Embedded Systems Concepts



Da-Wei Chang
CSIE, NCKU

*Code and Pictures on the slides are from
MicroC/OS-II, The Real-Time Kernel (2nd Edition), by Jean J. Labrosse*



Real-time Systems

- Cares about
 - Logical correctness
 - **Timing** correctness
 - deadline
- Predictable performance
- Two types of real-time (RT) systems
 - Soft vs. Hard (*described in the next slide*)



Soft RT vs. Hard RT

■ Soft RT

- Tasks are performed as fast as possible

■ Hard RT

- Must not miss the deadline
- Severe consequences if deadline misses...



Embedded Systems

■ Embedded system

- A computer built into a system and not seen as being a computer

■ Examples

- Internet of **Things**
- Robots
- Automotive
- Smart glasses
- Drones
-




Embedded Systems

- Three types of embedded software architecture
 - Foreground/background
 - Non-preemptive RTOS (kernel)
 - Preemptive RTOS (kernel)

Foreground/Background System

- Also called as *super loops*
 - A main loop + n ISRs, where $n \geq 0$
 - ISR = Interrupt Service Routine
- Used in low complexity, small systems





Tasks

■ Tasks

- Sometimes called as *threads*
- Programs in execution
 - A sequence of codes in execution
- Each task thinks that it owns the CPU



Foreground/Background System

■ Background

- Application with an infinite loop
- Task level, a single task

■ Foreground

- Interrupt service routine (ISR)
- Interrupt level
- Handle **asynchronous** events
- Used to perform critical operations

No RTOS in such a system!



Foreground/Background System

- The execution time of the application loop
 - Depends on the trigger situations of the interrupts and the ISR execution time
 - **Non-deterministic**
- A lot of micro-controller based systems are foreground/background systems



Non-preemptive or Preemptive RTOS/Kernel

- Responsible for managing tasks
- Ease application design
 - Provide services to allow creating multiple tasks
 - Provide services to allow schedule multiple tasks
 - ...
- Has kernel overhead
 - Time overhead
 - Depends on how often you invoke the kernel services
 - Space overhead
 - Depends on what is included in the kernel



RTOS Kernels

■ Real Time and Embedded Kernels

- FreeRTOS
- MicroC/OS-II, MicroC/OS-III
- Cs/OS2, Cs/OS3 (Cesium RTOS)
- VxWorks
- QNX
- RT-Thread
- ThreadX
- Zephyr
- RIOT
- Mbed
- ...



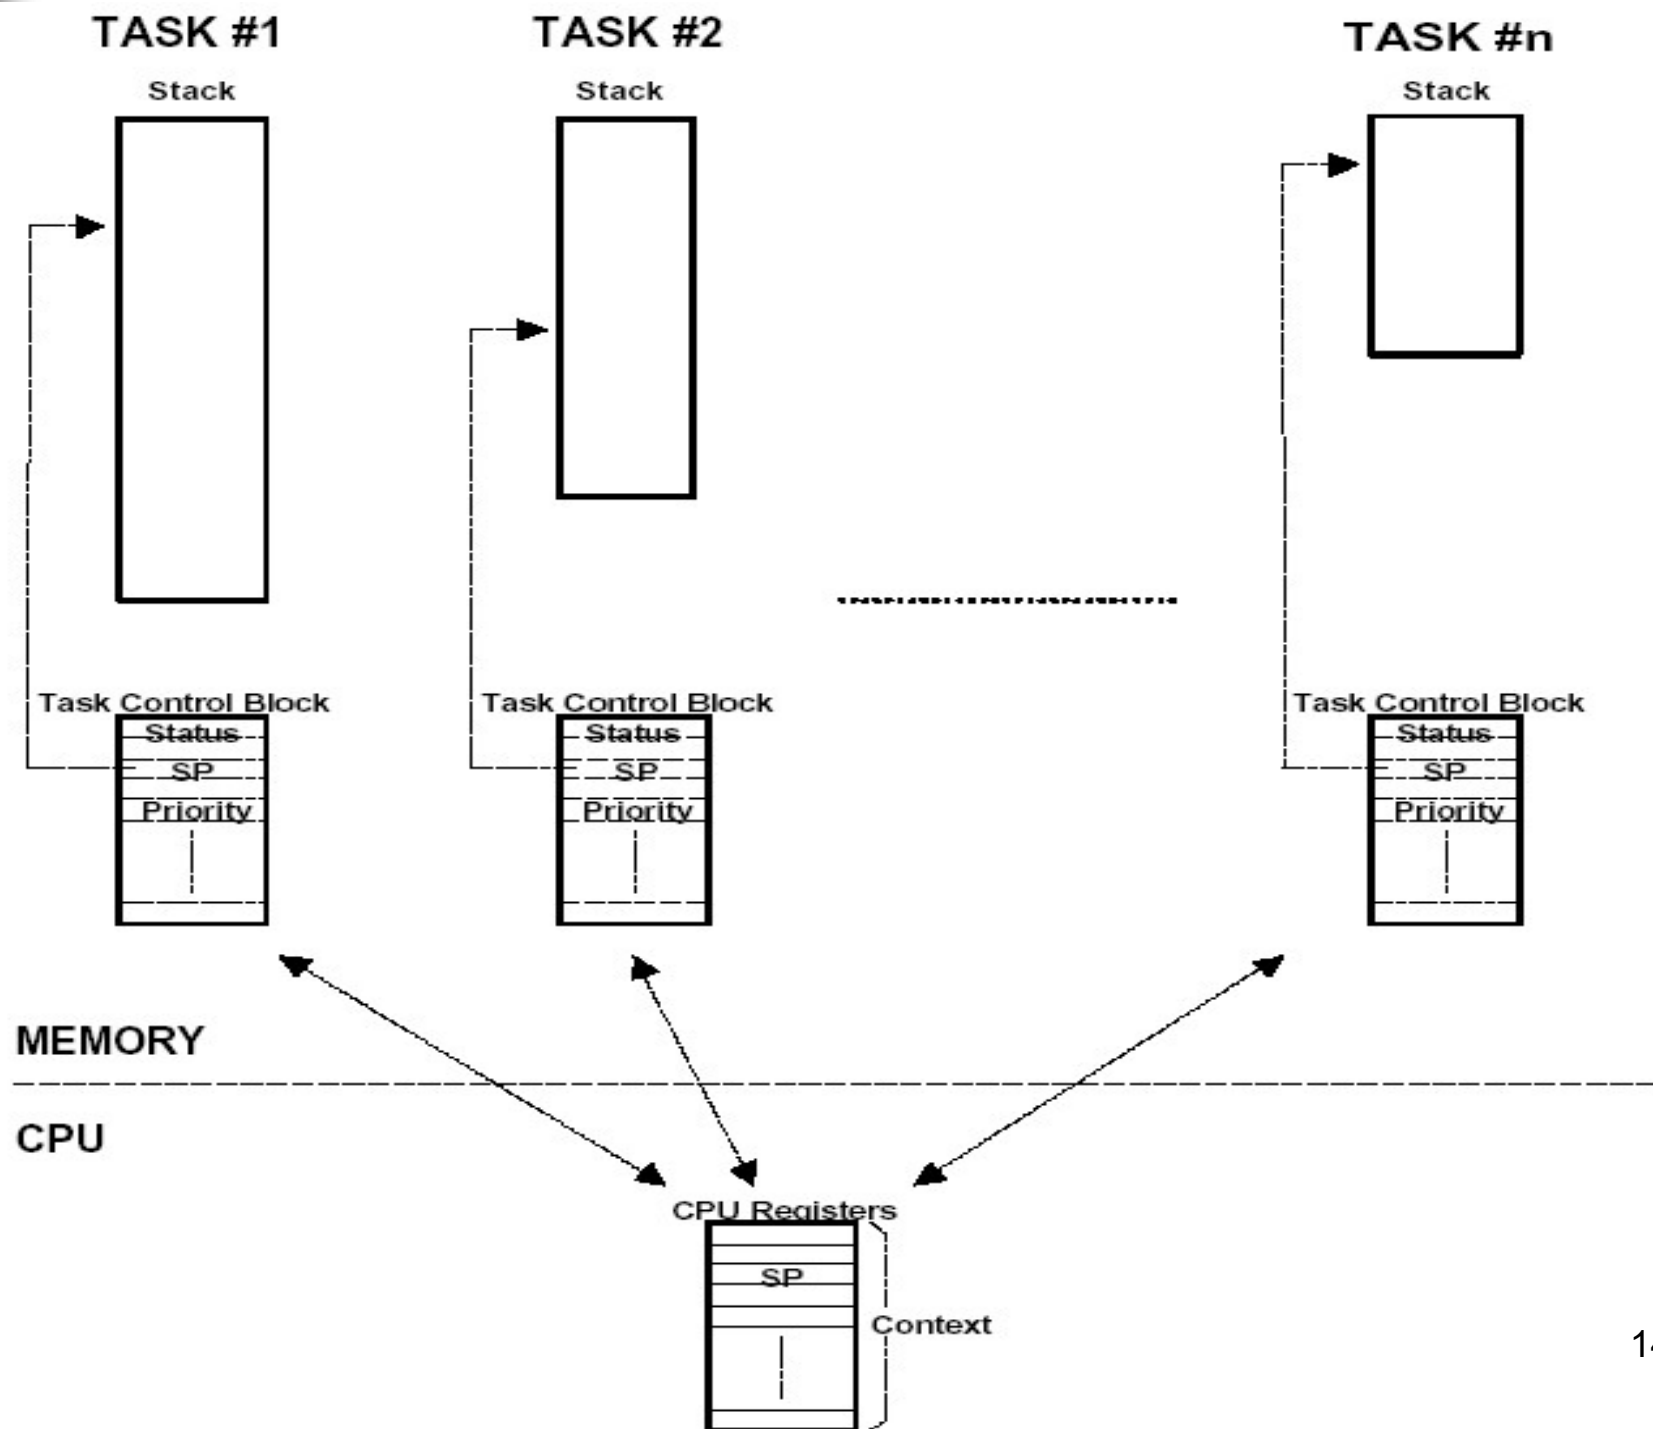
MultiTasking

- Running multiple tasks concurrently
 - Dynamically switching one task to another
- Benefits
 - Maximize the utilization of the CPU
 - Provide for modular construction of applications
 - Splitting application jobs into multiple tasks → applications are easier to design and maintain
- Each task thinks that it owns the CPU



Tasks

- Each task has the following
 - Priority
 - Stack
 - State
 - Space for storing values of CPU registers
 - For some systems, the space may be in the task's stack

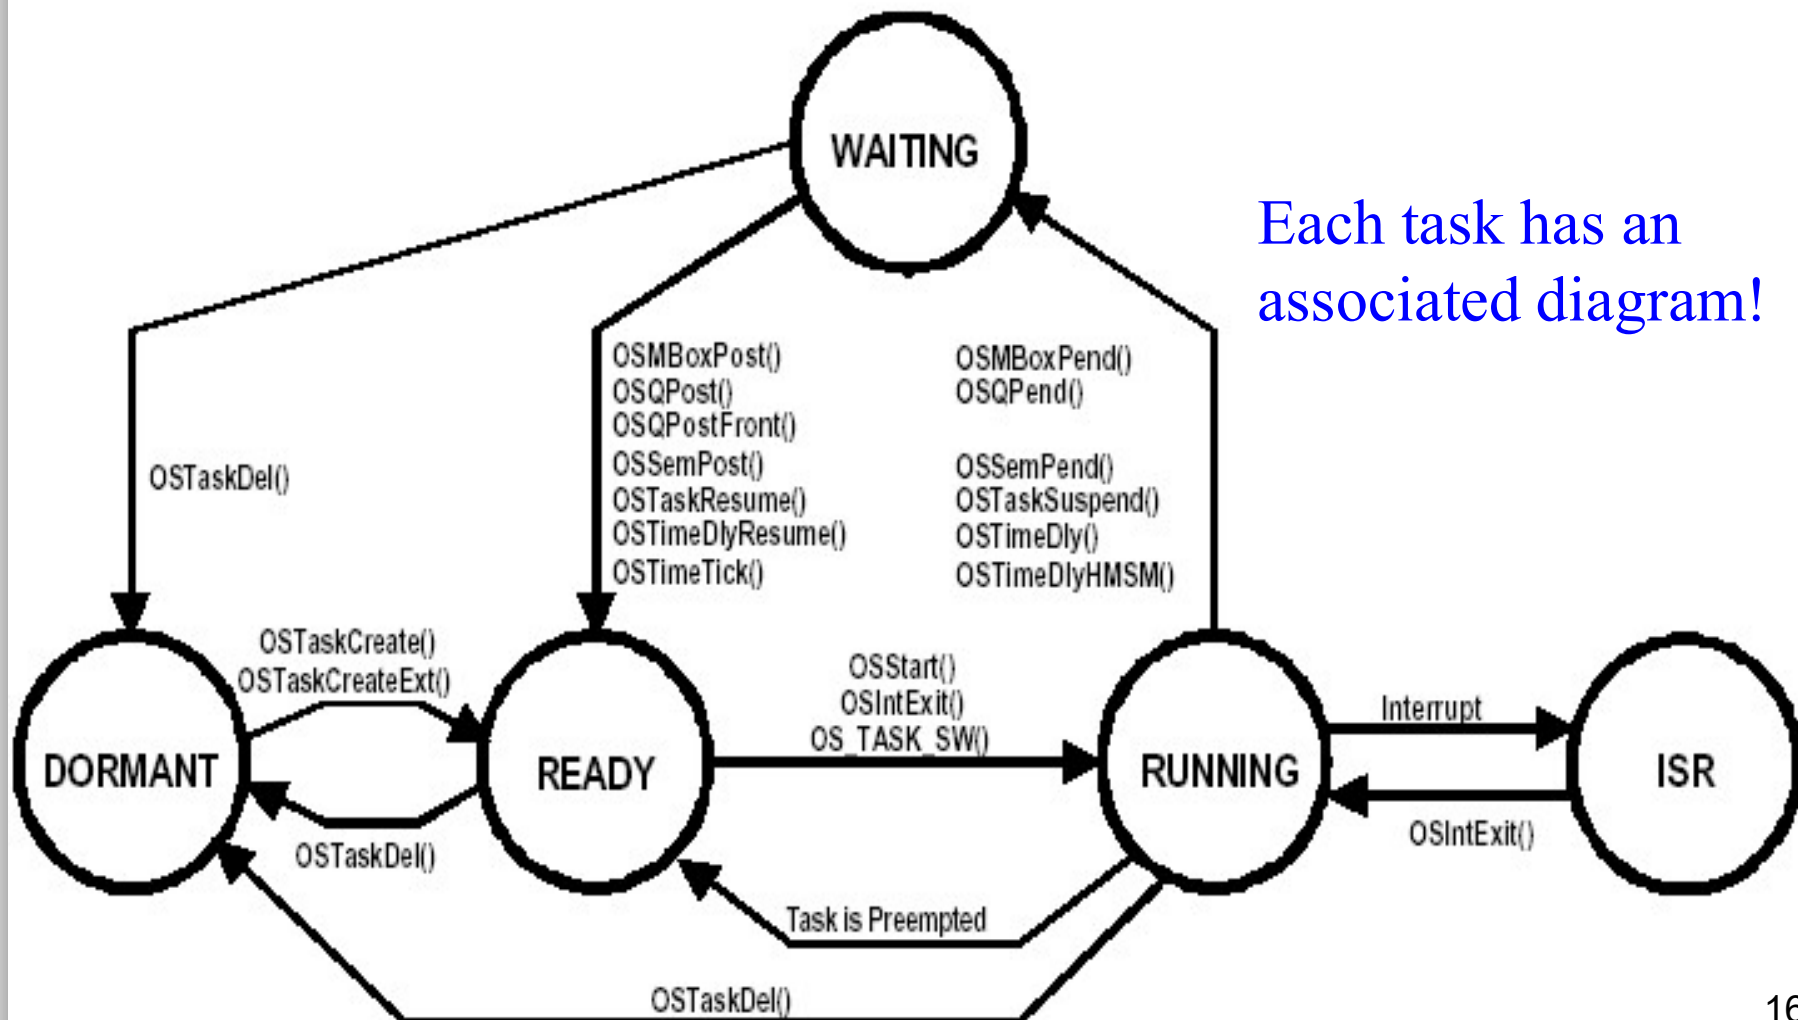




Context Switches

- Also called *task switches*
- Switch the CPU from one task to another
- Context switch jobs (from task A to task B)
 - Save context in A's stack
 - Restore context from B's stack
- Overhead is related to
 - Context size
 - Switching frequency

Task States





Schedulers

- A part of OS kernel
- Determines which task to run next
- Usually based on **priority**
 - The highest-priority task gets the CPU
 - In many cases, *assigning priority* is the job of the application designers.
- Priority-based scheduling algorithm
 - static priority vs. dynamic priority
 - For same-priority tasks
 - First-In-First-Out (FIFO), Round-Robin (RR)
- Other scheduling algorithms
 - Rate monotonic (RM), Earliest-Deadline-First (EDF)



Assigning Task Priorities

- In a general priority-based scheduling system, **assigning task priorities may not be an easy job** if you have many tasks
 - Needs to be considered carefully



Schedulers

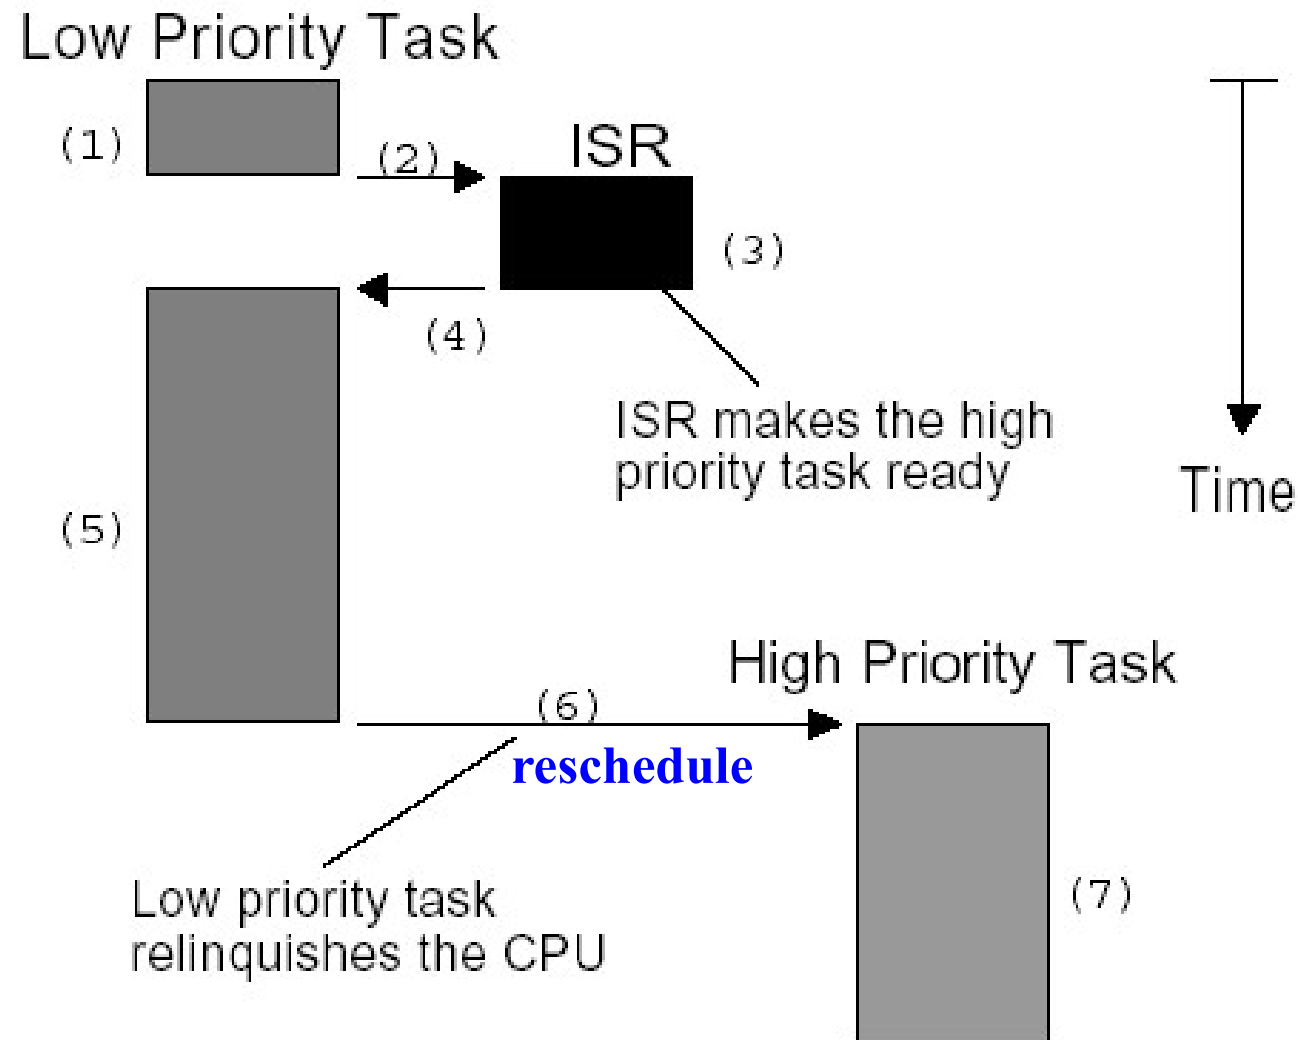
- When to perform re-scheduling?
 - Different in preemptive and non-preemptive kernels



Non-Preemptive Kernels

- Tasks should **explicitly** give up the CPU to allow rescheduling
 - Cooperative multi-tasking
 - e.g., Win95
 - Asynchronous events are still handled by ISRs
 - ISR can **preempt** current task
 - ISR can make a higher-priority task ready
 - **But**, ISR still returns to the **interrupted** task
 - *See the next slide*
- Less responsive
- + Less overhead to guard shared resources

Non-Preemptive Kernels

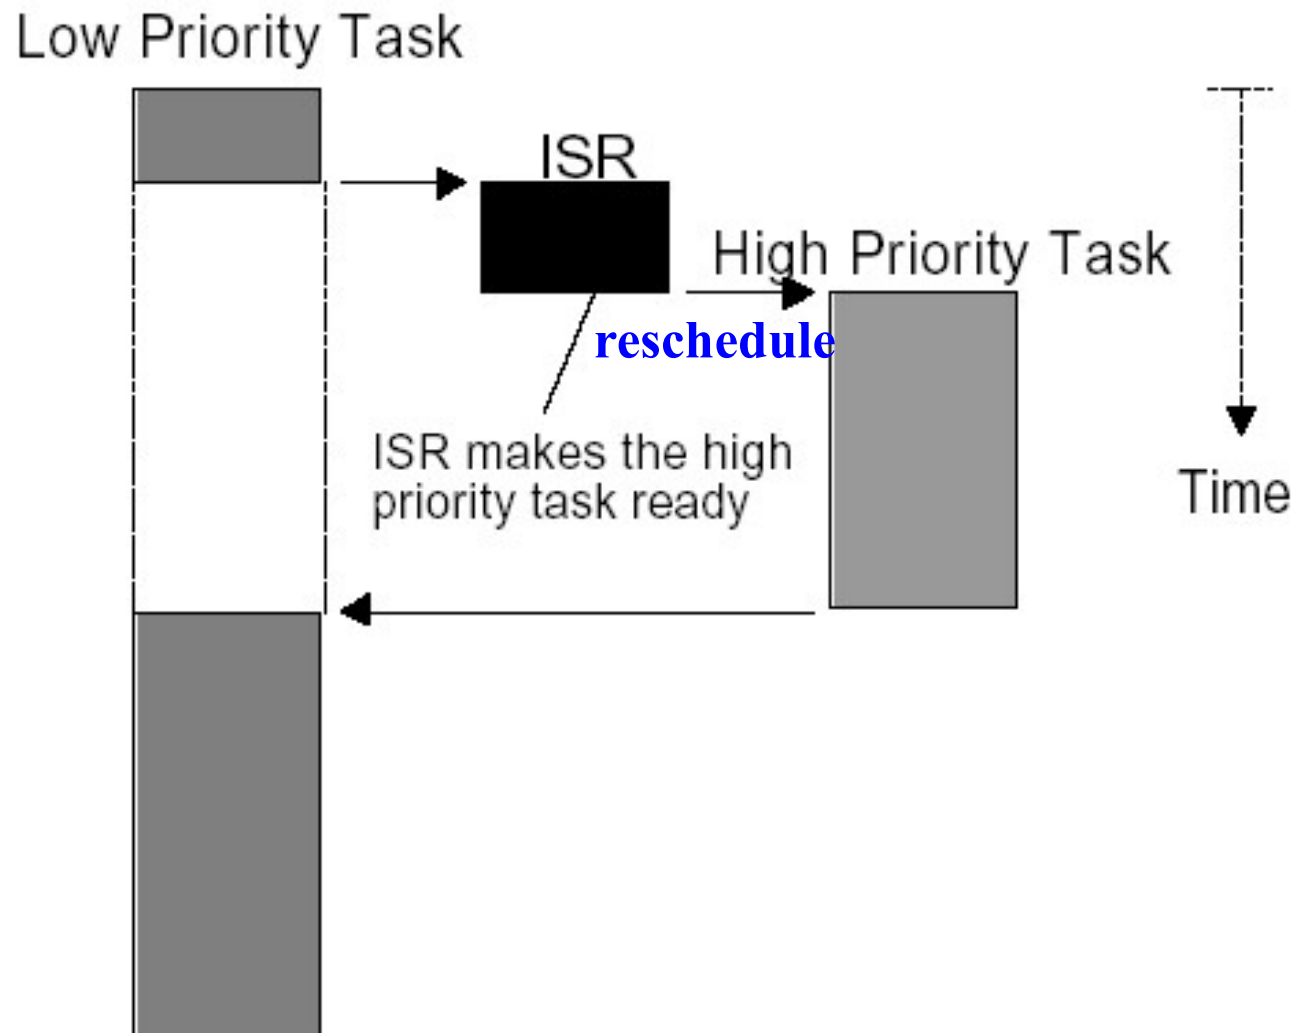




Preemptive Kernels

- Tasks are forced to give up CPU when needed
- Whenever a higher priority task is ready, the current task is preempted
 - *See the next slide*
- + More responsive
- More overhead to guard shared resources
- Most RTOSs are preemptive kernels

Preemptive Kernels





Preemptive Kernels

- On preemptive kernels
 - it is easier to **determine** when the highest priority task can get the CPU



Reentrant vs. Non-Reentrant Functions

```
void strcpy(char *dest, char *src)
{
    while (*dest++ = *src++) {
        ;
    }
    *dest = NUL;
}
```

Reentrant

```
int Temp;

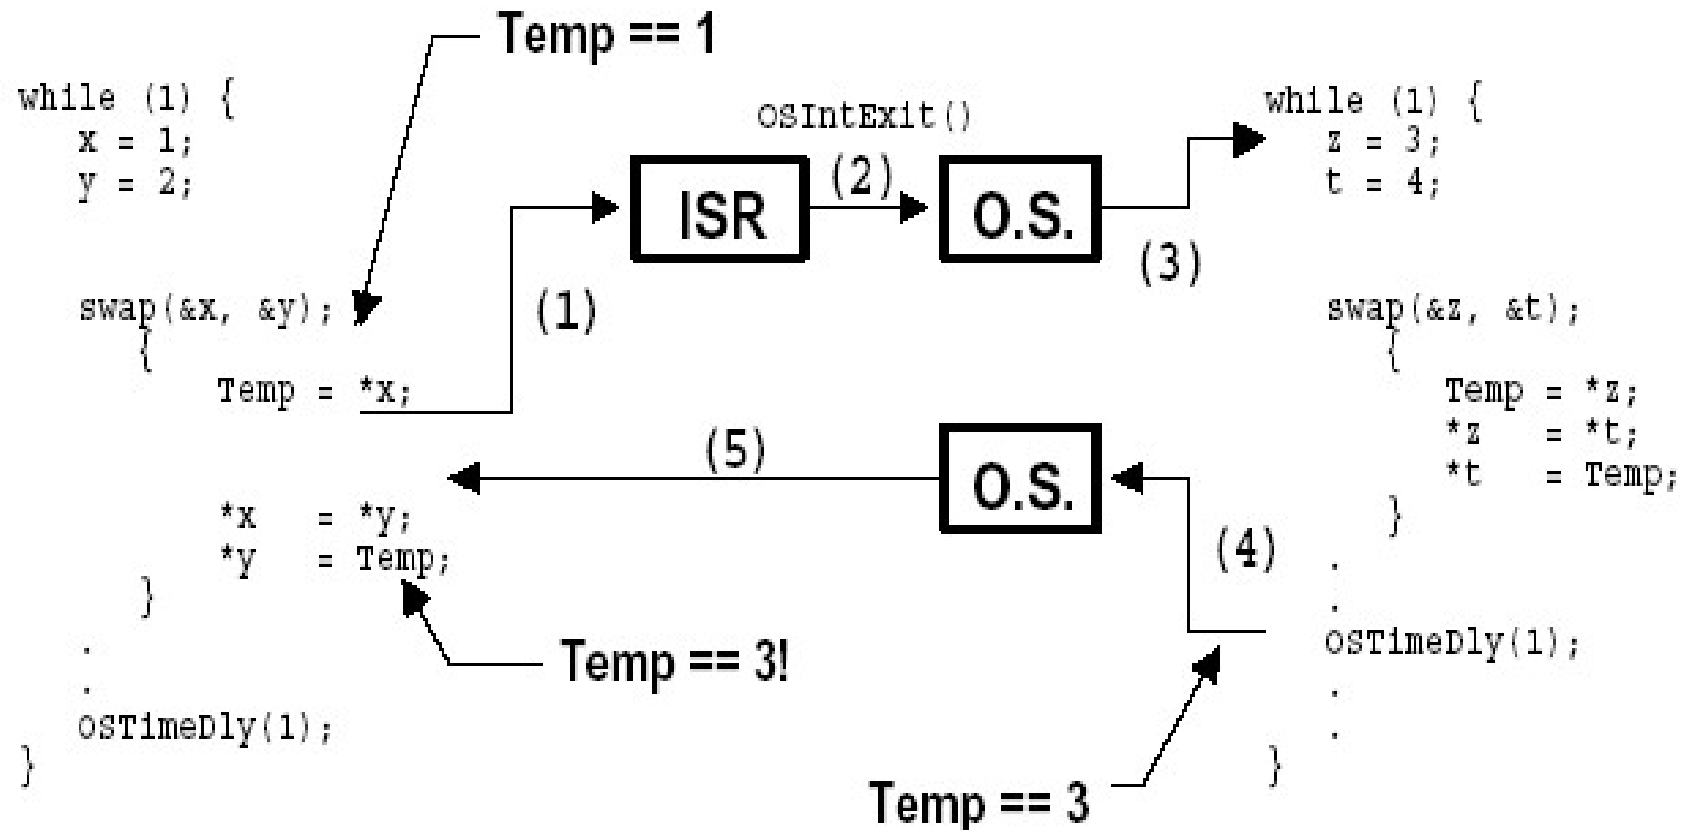
void swap(int *x, int *y)
{
    Temp = *x;
    *x = *y;
    *y = Temp;
}
```

Non-Reentrant

An Example of Non-Reentrant Function

LOW PRIORITY TASK

HIGH PRIORITY TASK





Making a Function Reentrant?

- Use local variables only



Resources

- Any entity used by a Task/ISR
 - Physical resources
 - e.g., IO devices like printer, keyboard,...
 - Logical resources
 - e.g., data structures like variables, arrays, structures...



Resources

■ Shared resources

- Resources used/accessed by more than one tasks/ISRs

■ **Mutual exclusion**

- At any given time, only one task can access the (shared) resource



Critical Sections

- A sequence of code that accesses one or more shared resources
- Ensure *mutual exclusion*
- Approaches for ensuring mutual exclusion
 - Disable interrupts
 - Use test-and-set operations
 - Disable scheduling
 - Use semaphores



Mutual Exclusion

- Disable interrupts
 - `OS_ENTER_CRITICAL()` and `OS_EXIT_CRITICAL()` in uC/OS-II
 - Issues
 - Which interrupts?
 - Which interrupt levels?
 - Save interrupt state before disabling?



Mutual Exclusion

- Test-and-set operations
 - **Test** a global variable
 - 0 → access resource, 1: don't access
 - Have to **Set** the variable when testing its value
- Implementation method 1
 - test-and-set or similar hardware instructions
- Implementation method 2
 - Disable interrupt for accessing the variable



Mutual Exclusion

■ Disable Scheduler

- *Less performance impact* than disabling interrupts
- Rescheduling **does not** occur even if a higher priority task is made ready by an ISR
 - ISR still returns to the interrupted task
 - Similar to non-preemptive kernel
- Works if the task does not share resources with ISRs
- In uC/OS-II, the related API are
 - OSSchedLock()/OSSchedUnLock()



Mutual Exclusion

■ Semaphores

– Used to

- Control access to a shared resource
- Signal the occurrence of an event
- Task synchronization

■ Binary vs. Counting

■ For Mutual Exclusion


- Acquire the semaphore (i.e., key) before accessing the resource
- Release the semaphore after accessing the resource



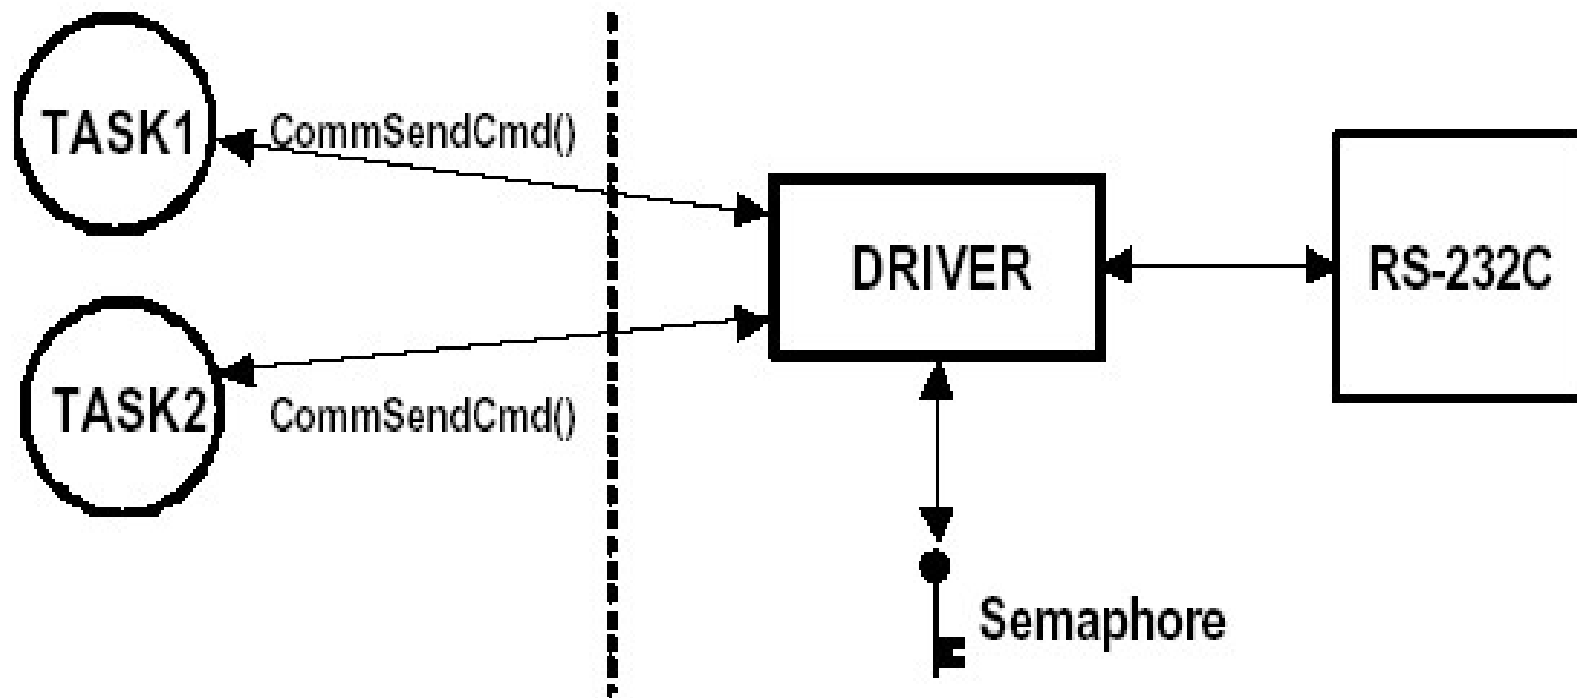
Semaphores

- Operations
 - **Wait** (acquire)
 - **Signal** (release)
- Wait may cause suspension/sleep
 - When waiting a semaphore with value 0
- Signal may cause some tasks to be resumed/wake-up
 - Wake up the first suspended task
 - Wake up the highest-priority task ←uC/OS-II

A Semaphore Example

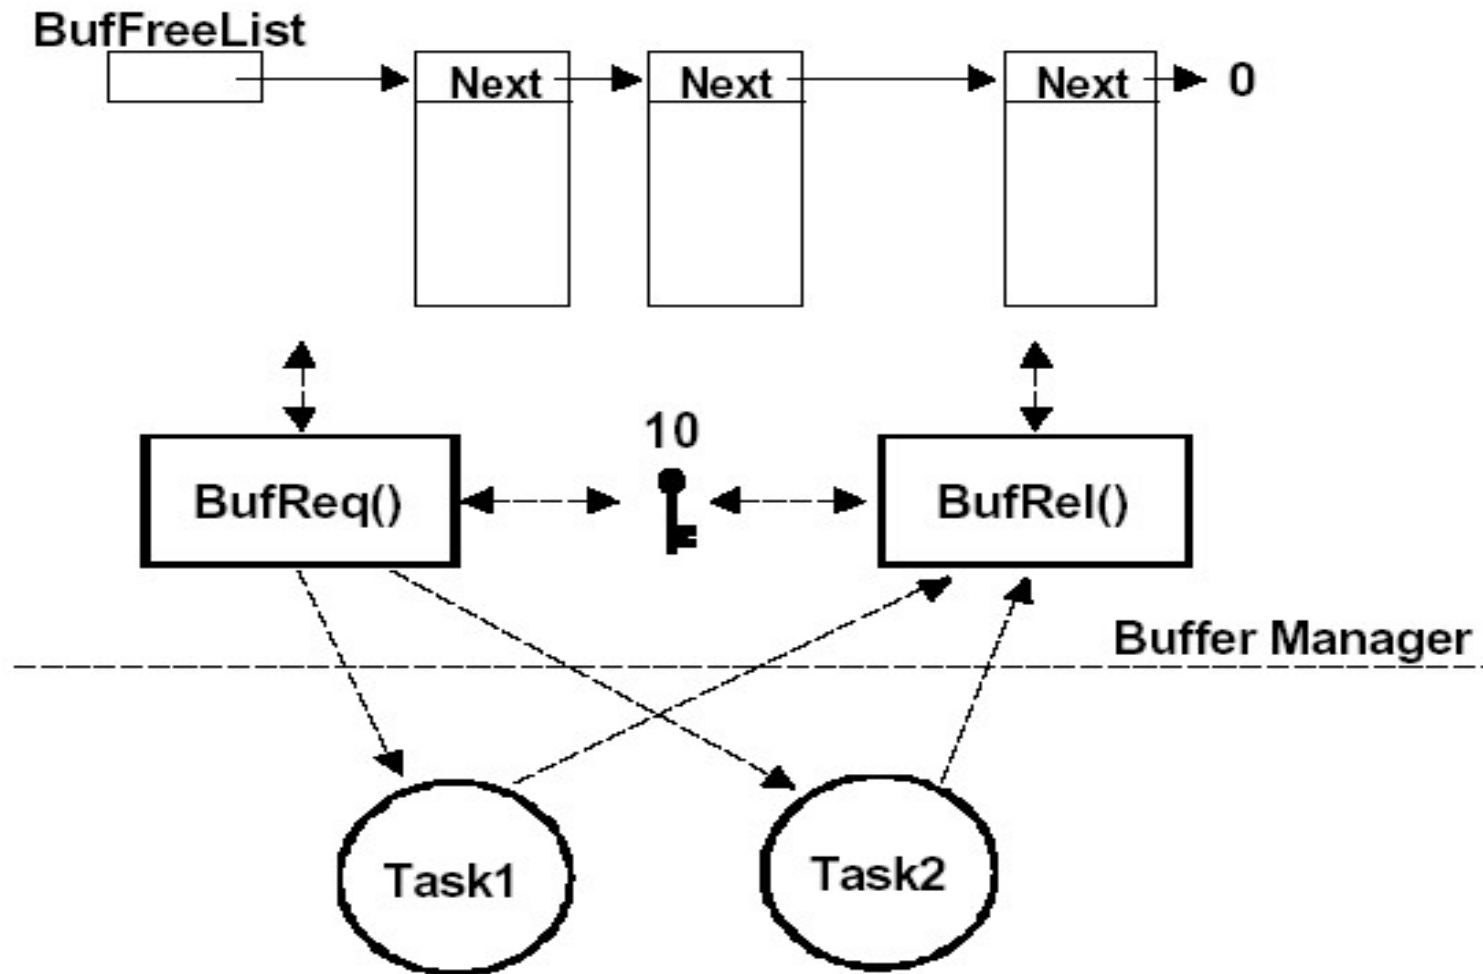


Encapsulating Semaphore Operations



Semaphore operations are encapsulated in CommSendcmd()₃₇

Another Semaphore Example



Counting semaphore with value = 10 (i.e., 10 keys)

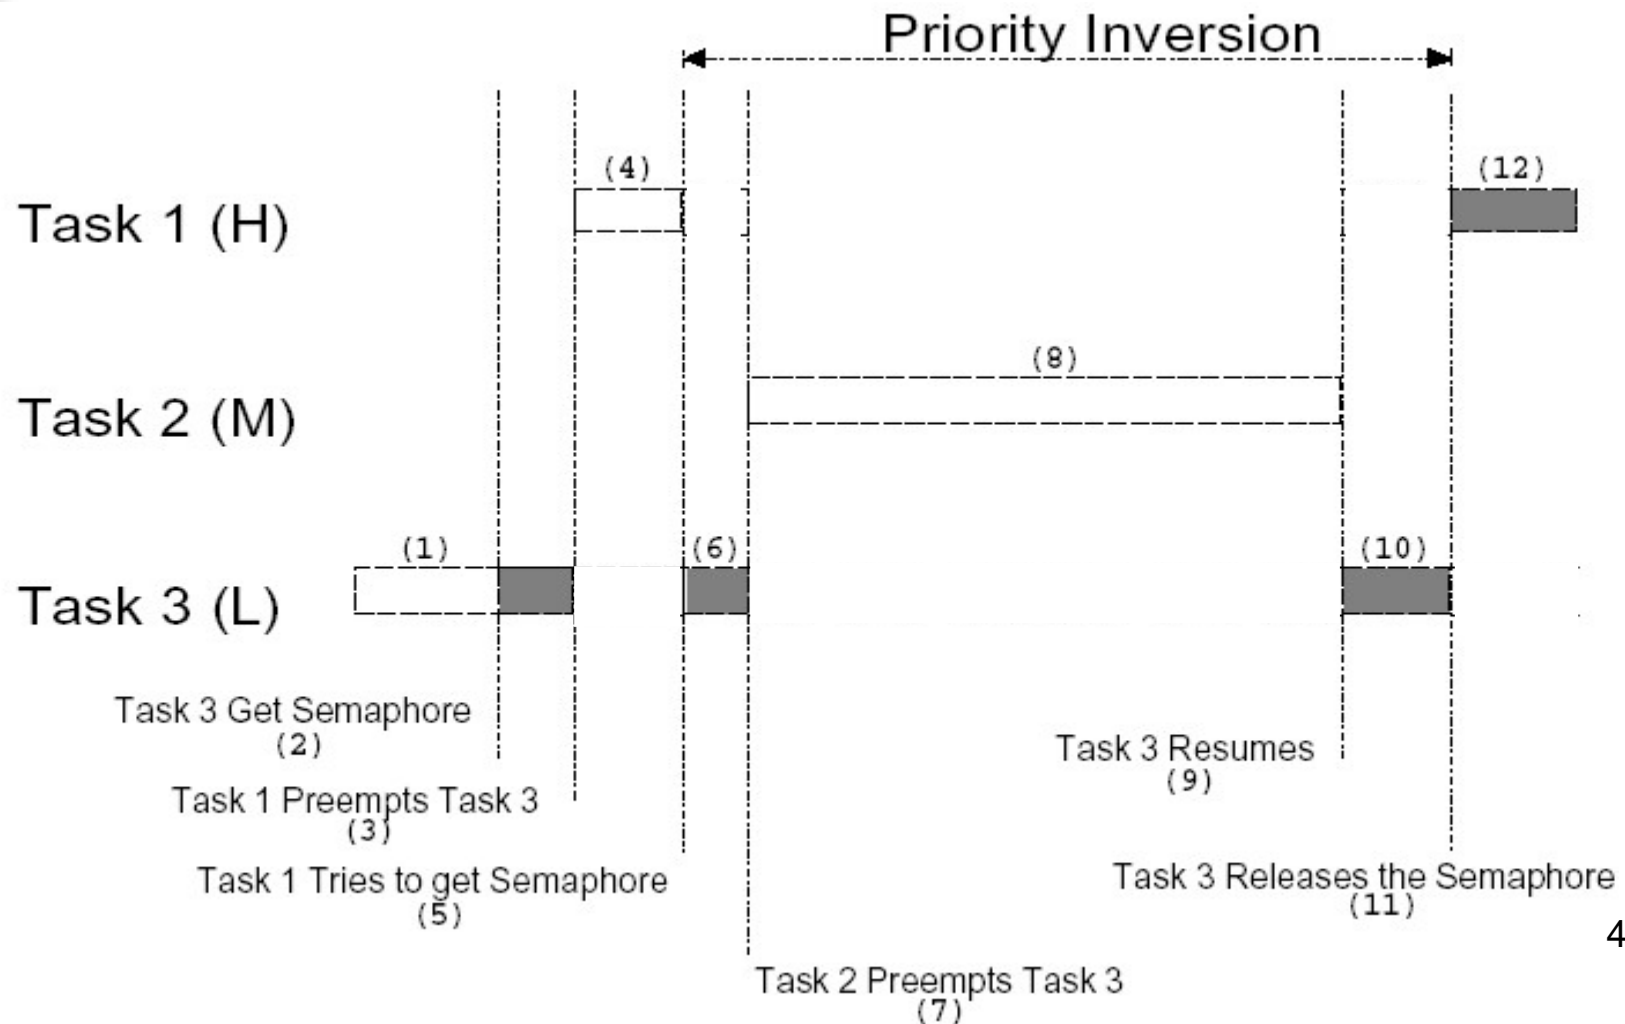


Deadlock


- Two or more tasks wait for resources held by the others
 - Causing system hangs
- Prevent deadlock
 - Acquire all resources before proceeding
 - Acquire resources in the same order
- Breaking deadlock
 - Timeout-based waiting
 - Widely supported in RTOSes!

Priority Inversions

- A problem in RT systems



Priority Inheritance Protocol (PIP)



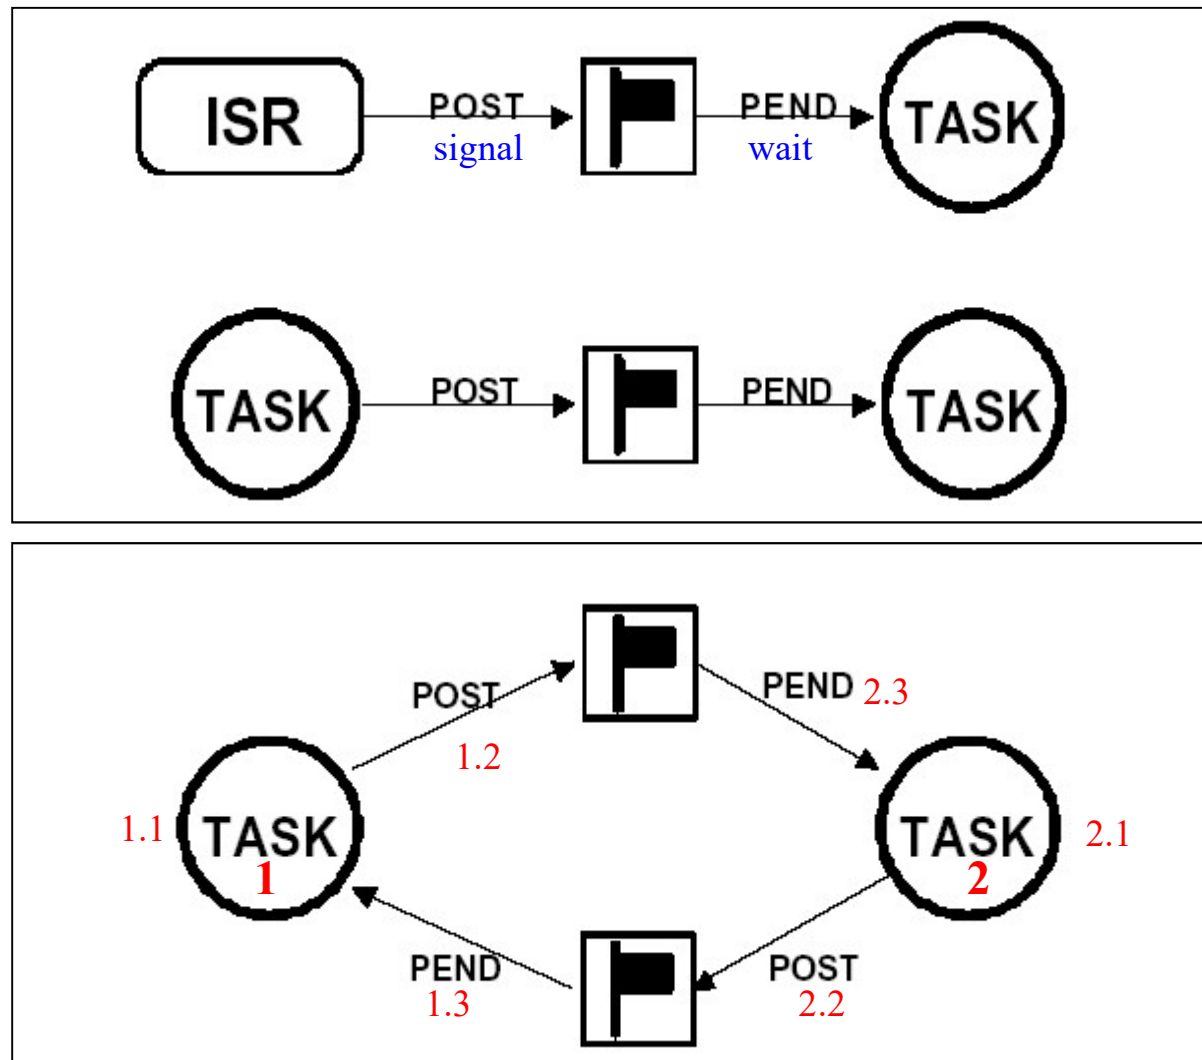
Some level of priority inversion can not be avoided!!



Synchronization

- Task-Task, Task-ISR
- Wait for an event to occur
- Approaches
 - Semaphores
 - Event flags

Synchronization -- Semaphore



Synchronization -- Semaphore


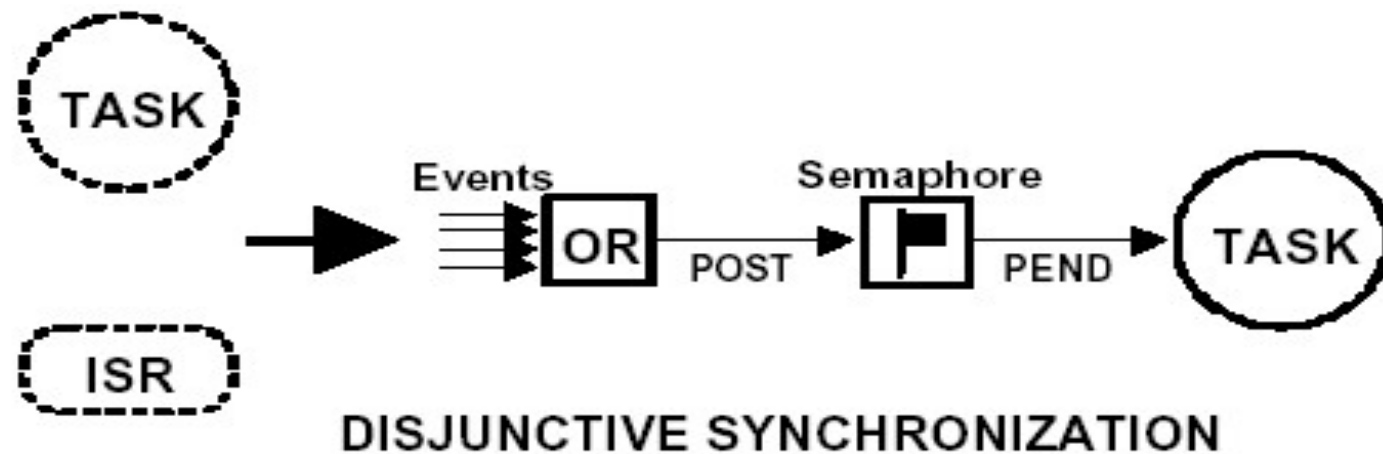
```
Task1()  
{  
    for (;;) {  
        Perform operation;  
        Signal task #2;                (1)  
        Wait for signal from task #2;  (2)  
        Continue operation;  
    }  
}  
  
Task2()  
{  
    for (;;) {  
        Perform operation;  
        Signal task #1;                (3)  
        Wait for signal from task #1;  (4)  
        Continue operation;  
    }  
}
```



Synchronization – Event Flags

- Synchronize with the occurrence of **multiple** events
- You can define different eventTypes
 - Disjunctive (OR)
 - Conjunctive (AND)

Synchronization – Event Flags

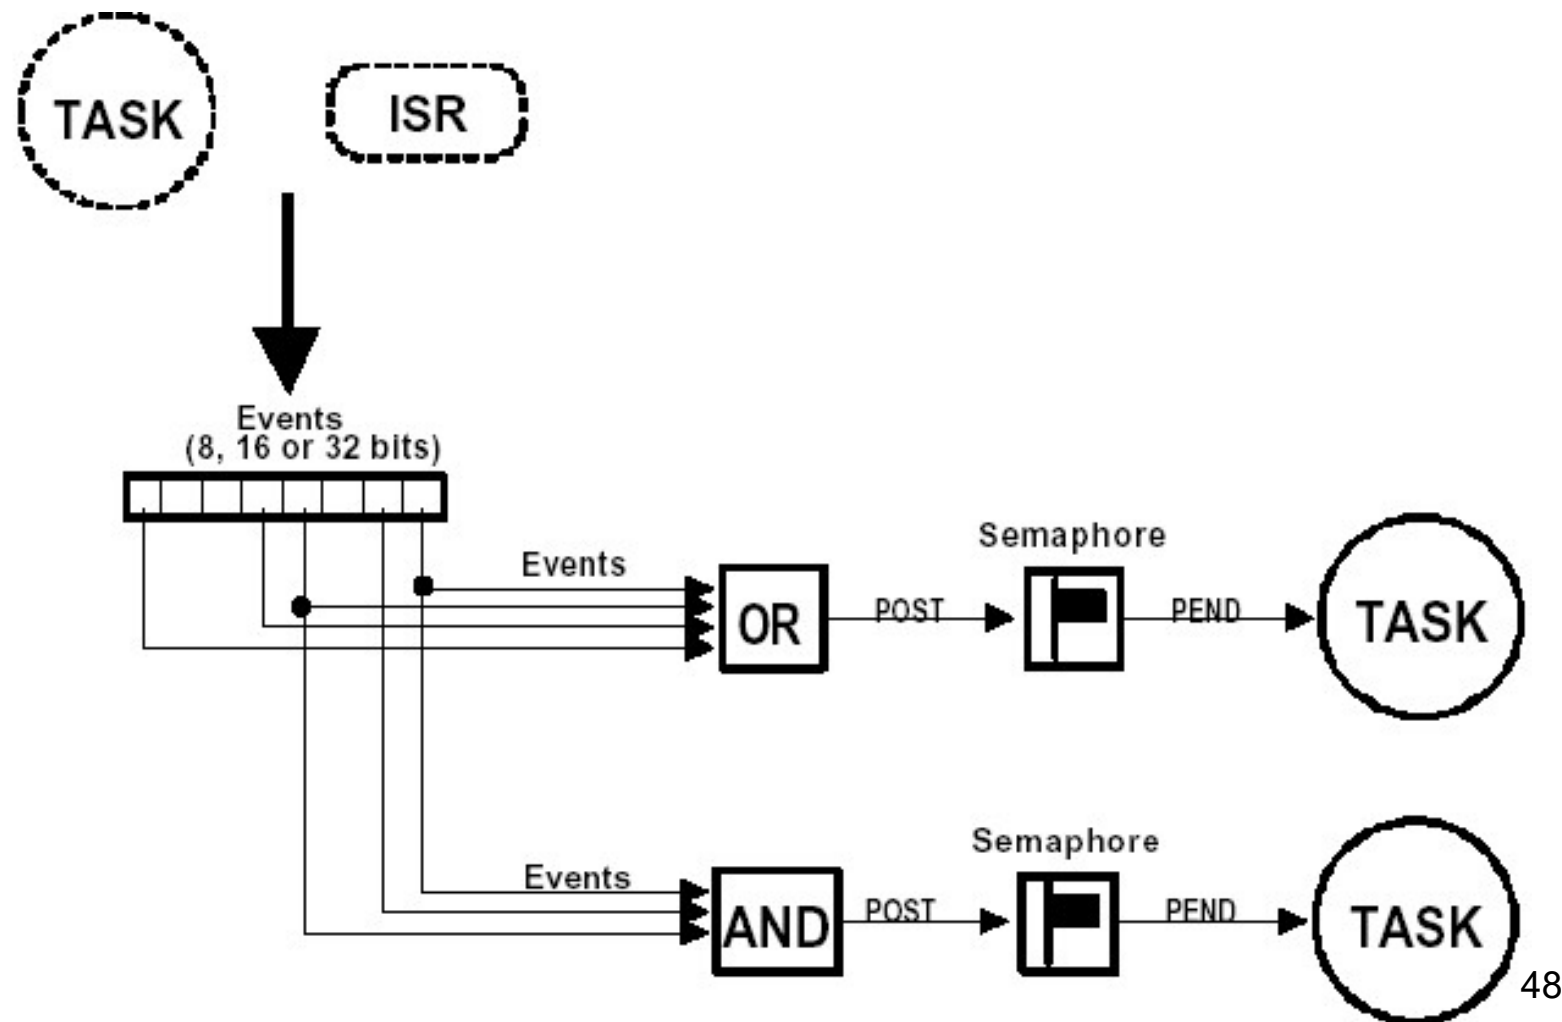




Synchronization – Event Flags

- Events are grouped
 - Usually represented as *bitmaps*
 - Task/ISR can set or clear events
 - *See next slide*

Synchronization – Event Flags





Inter-Task Communication

- For a task or an ISR to exchange information with another task
- Approaches
 - Global variables
 - Must ensure exclusive access
 - Mailboxes/Message queues
 - The most widely used inter-task communication mechanism in an RTOS



Mailbox/Message Queues

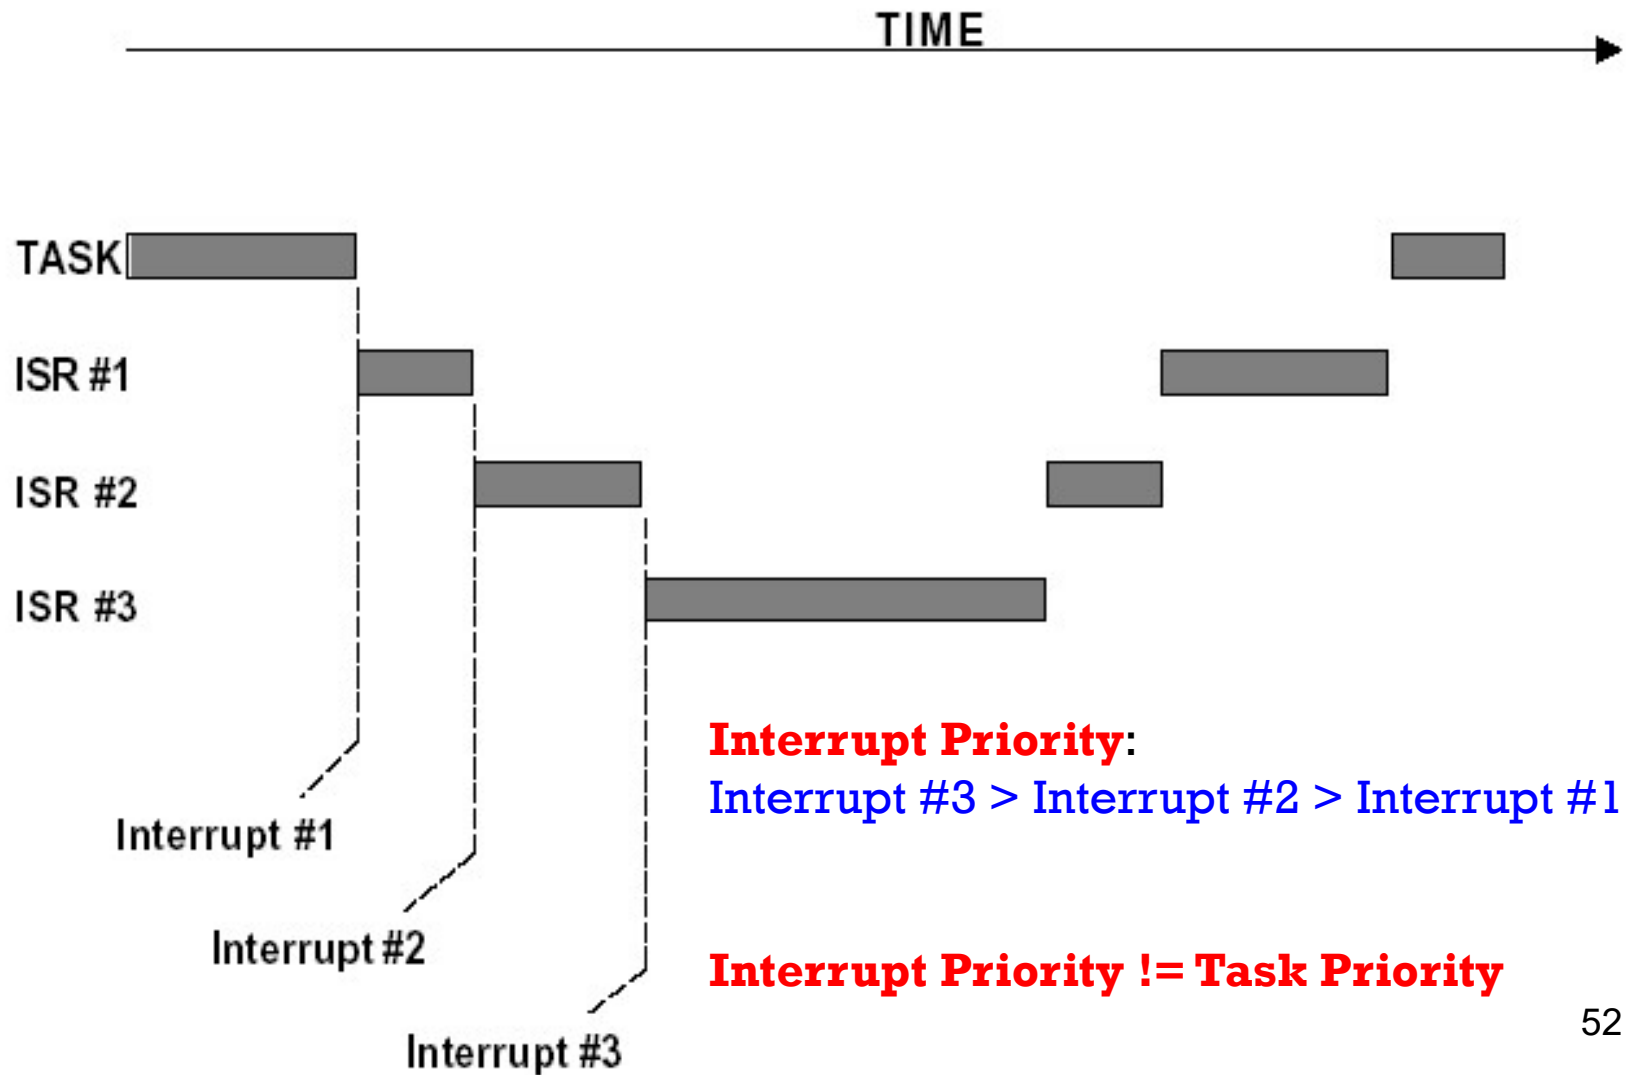
- A message-exchanging facility
- Can contain one or more messages
 - In uC/OS-II,
 - A mailbox can have only one message
 - A message queue can have more than one messages
- One or more tasks can receive on a mailbox
- Send message to a mailbox, not a specific task
 - Indirect communication
- A task may suspend when it wants to receive a message from an empty mailbox
 - Wakeup when a message comes



Interrupts

- Inform the CPU that an asynchronous event has occurred
- Improve CPU efficiency
 - Avoid **polling** devices
- After recognizing the interrupt, the CPU
 - Save (part of the) context
 - Jump to ISR (Interrupt Service Routine)
- When ISR is returned, control jumps back to
 - The interrupted task (non-preemptive-kernel)
 - The highest-priority task (preemptive-kernel)

Nested Interrupt





Interrupt Related Time

■ Interrupt Latency

- Between interrupt occurrence and the time to start executing the first ISR instruction

■ Interrupt Response

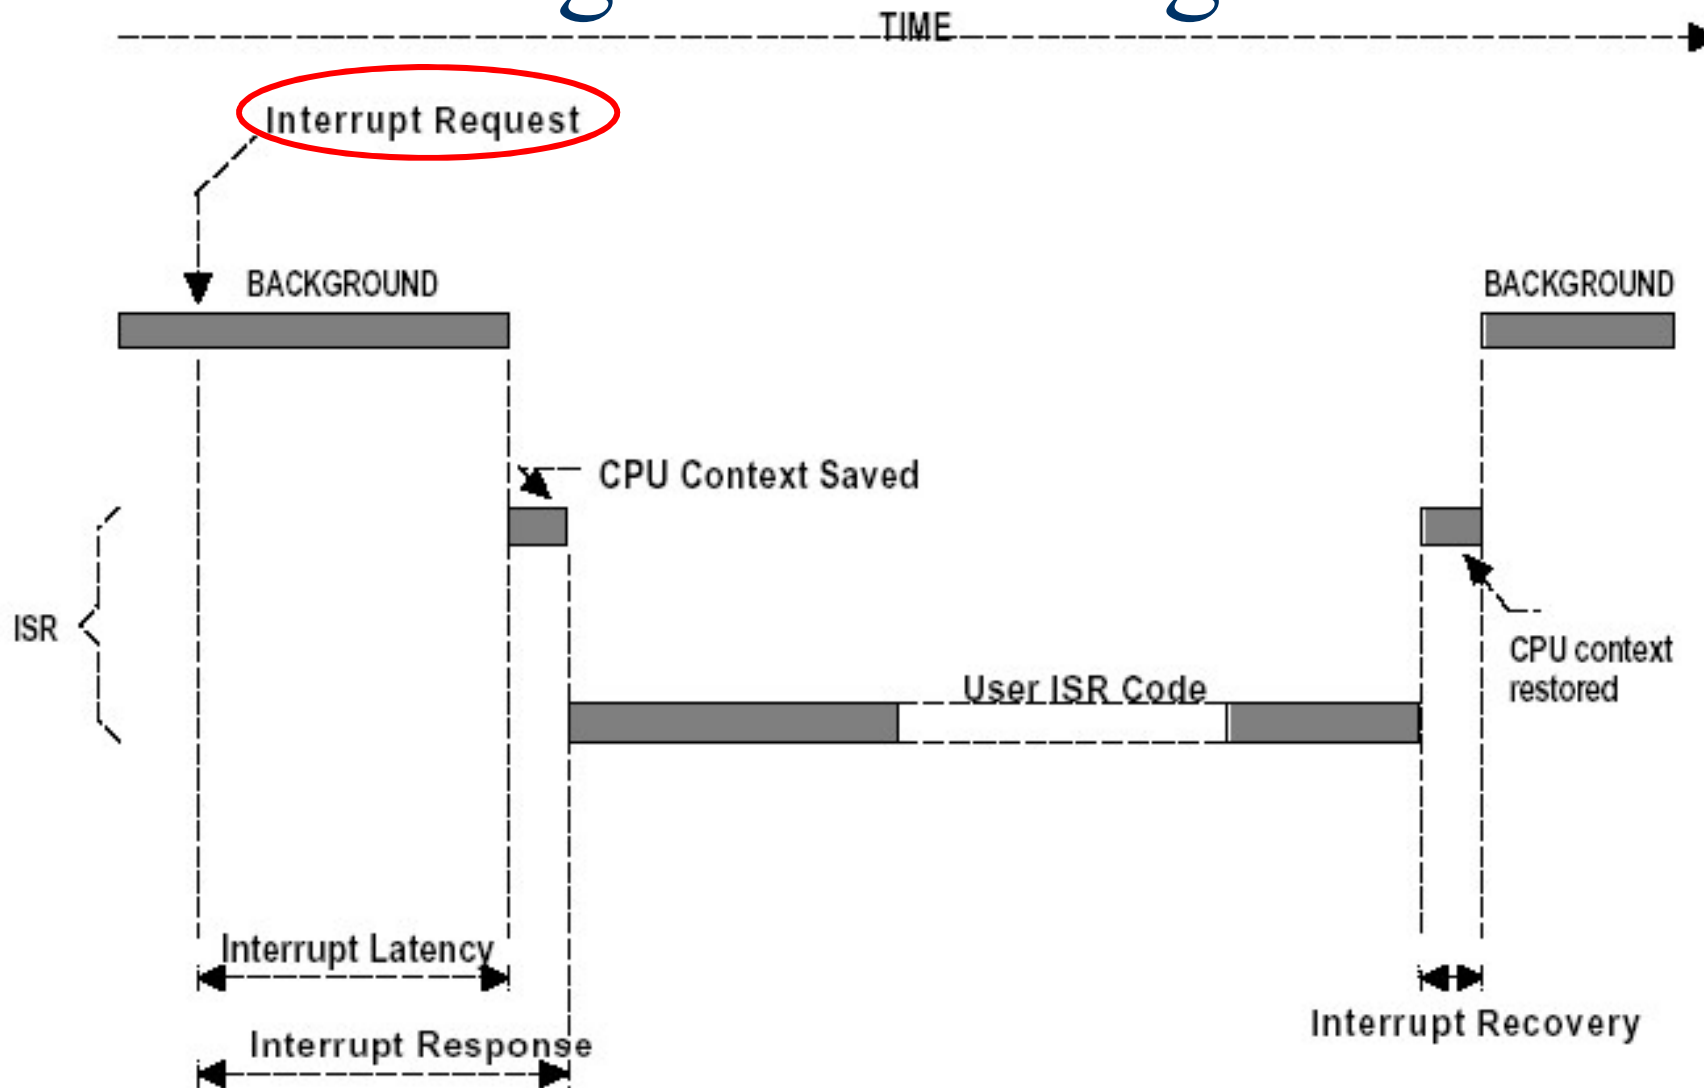
- Between interrupt occurrence and the start of the **user ISR code that handle the interrupt**

■ Interrupt Recovery

- Time required for the CPU to return to the interrupted/highest-priority task

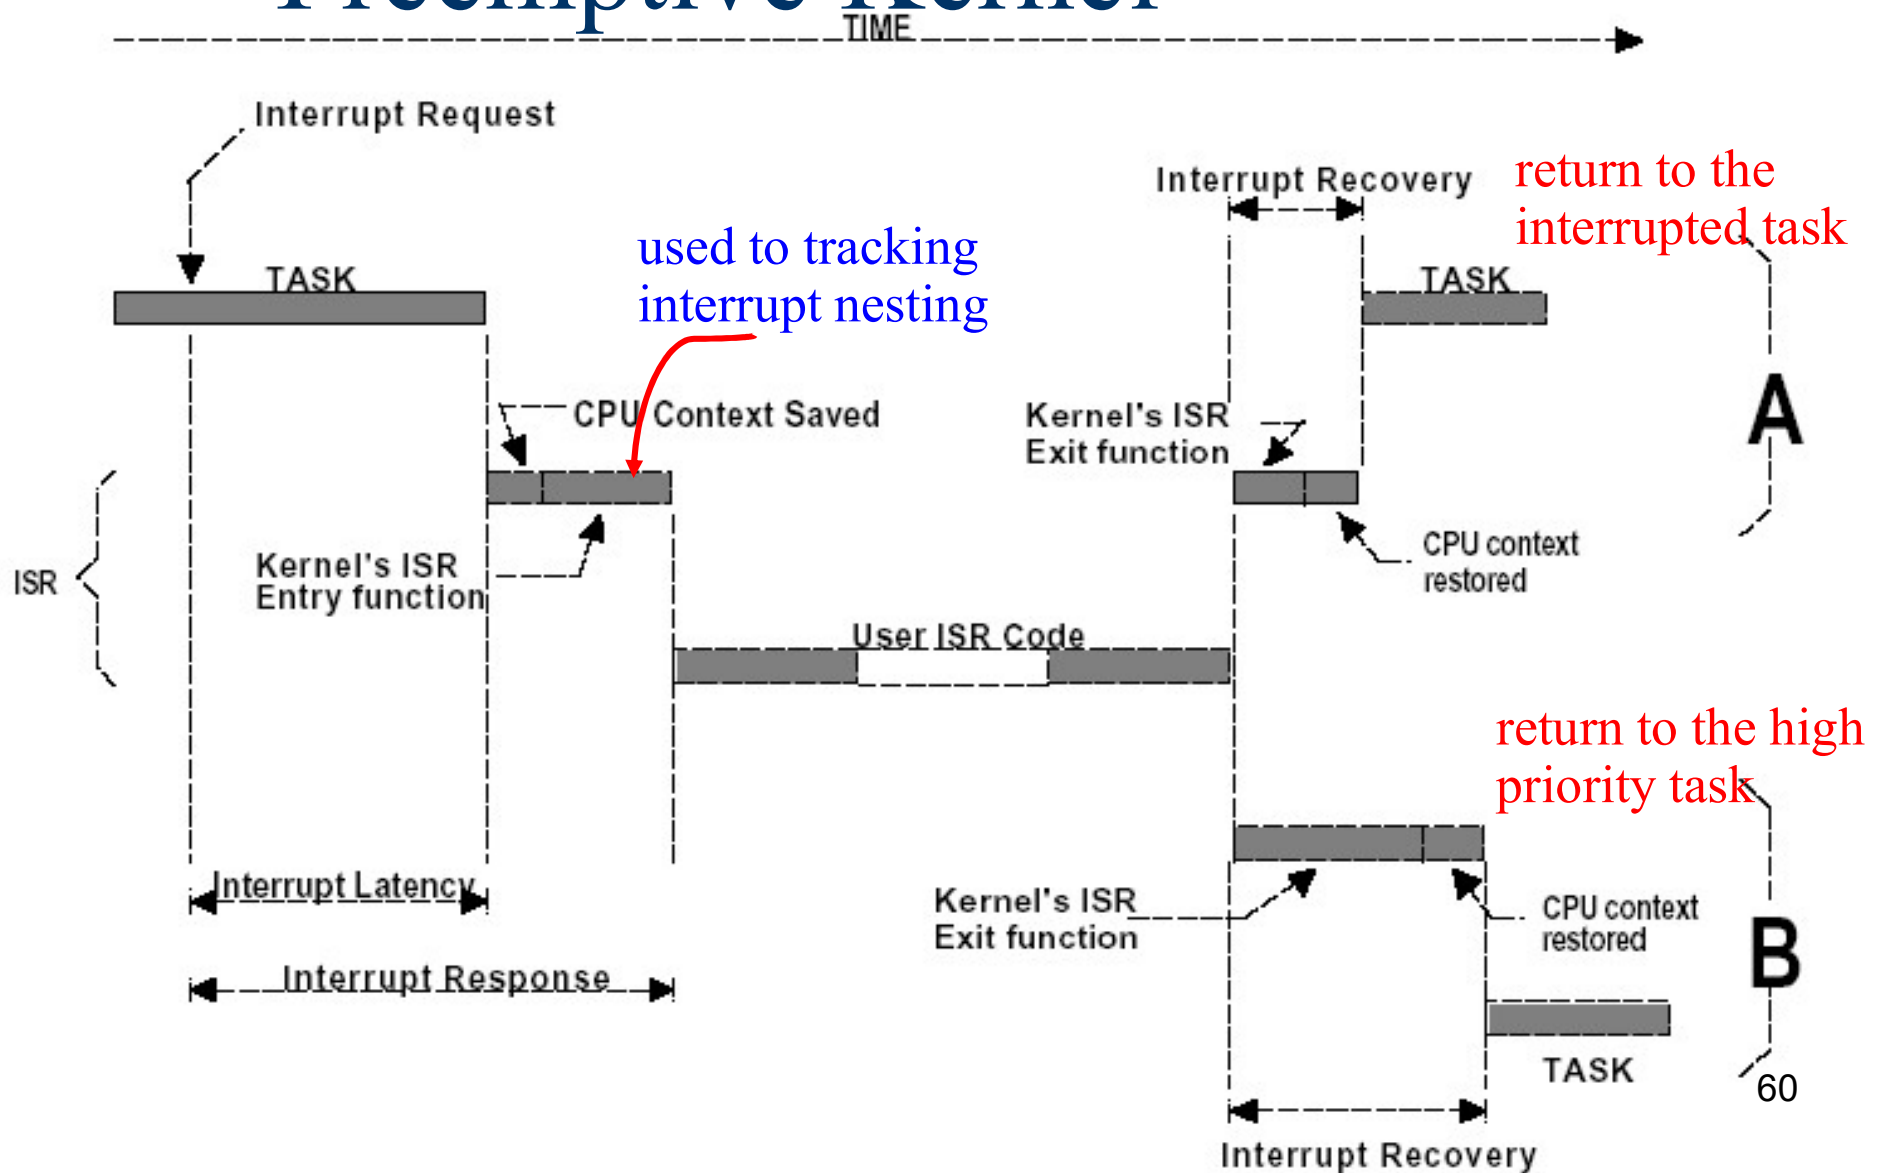
Interrupt-Related Time

-- Foreground/Background



Interrupt-Related Time

-- Preemptive Kernel





ISR Processing Time

- Interrupt occurs → jobs need to be done
- You can
 - Execute the jobs in the ISR, or
 - Execute the jobs in a task
 - e.g., ISR1/ISR2 in OSEK & AUTOSAR
- It depends!!
 - **Rule:** ISR should not consume too much time



Clock Tick

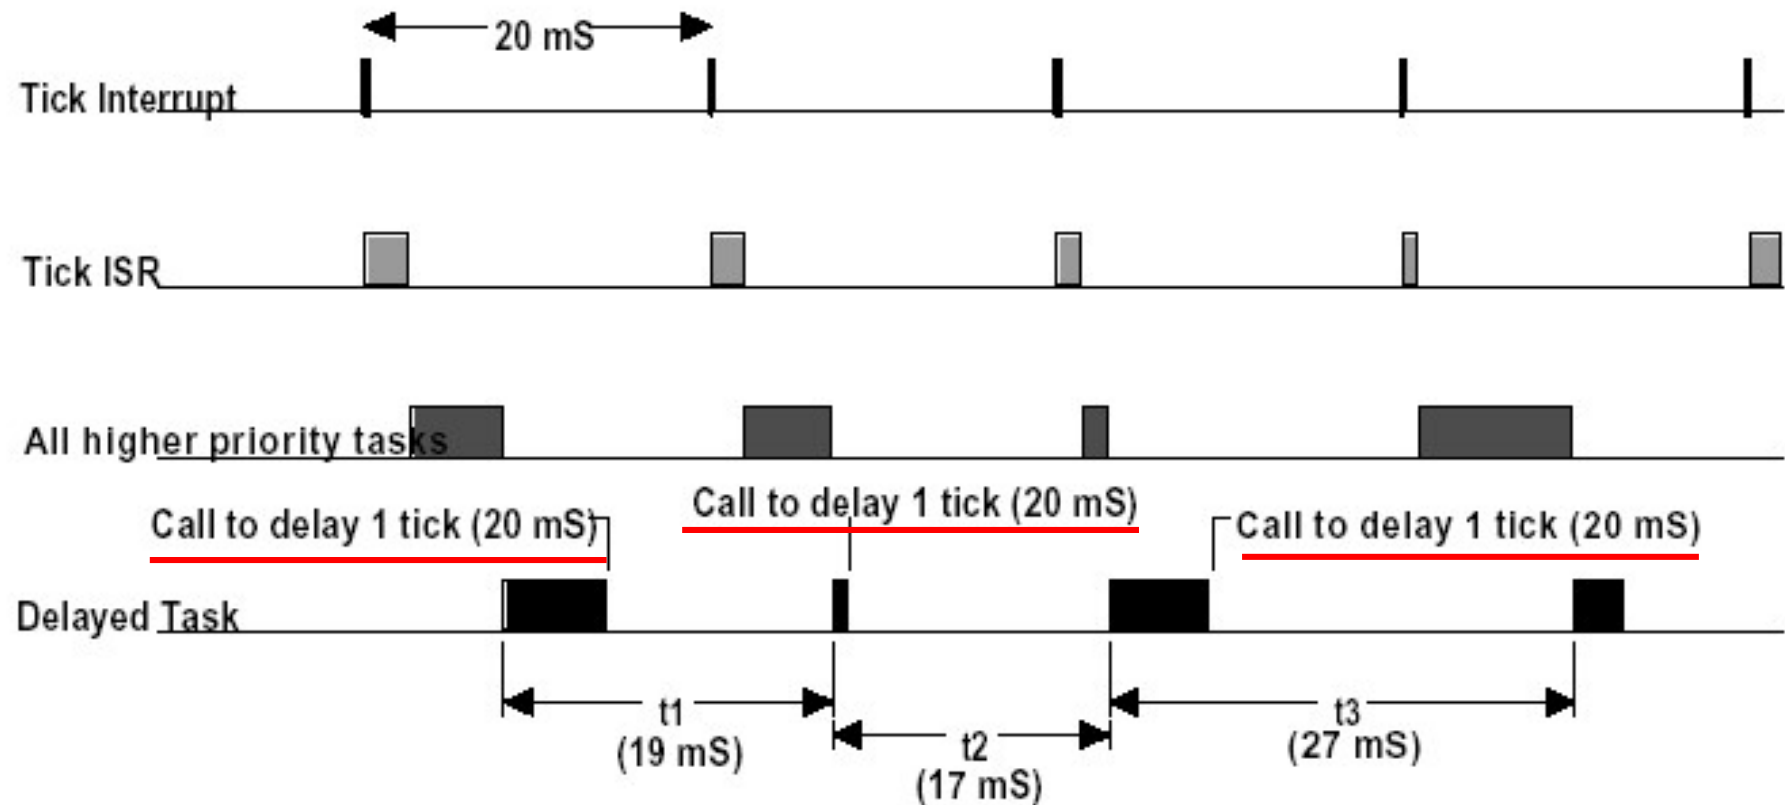
■ Timer interrupt

- A special interrupt that occurs periodically
- Can be used to measure time
- Can also be viewed as the system's heartbeat


■ Clock tick

- Time between successive time interrupts
- Typically, 1-200ms
- Allows
 - Task delay, timeout for waiting....
 - The basic delay time unit is a tick....
- Overhead vs. tick resolution

Delay Jitter – Case 1



Delay Jitter – Case 2





Delay Jitter

■ Possible Solutions

- Rearrange task priorities
- Write critical code in ISR

■ Delay jitter increases as the priority of the task decreases...



Memory Requirement

- *Memory size matters!*
- Memory Size Factors (when a multi-tasking kernel is used)
 - Application data/code size
 - Mainly depends on the application programmer
 - Use a high code-density ISA
 - Utilize compiler optimizations
 - Kernel data/code size
 - Data structures for managing the tasks...
 - Kernel code
 - Depends on the kernel functionality
 - Also depends on ISA and compiler



Memory Requirement

■ Memory Size Factors

- Task stack size

- Contains local variables, arguments
- Equal for all tasks?
- Consider
 - Too many arguments?
 - Use large local arrays, structures....?
 - Too deep function/interrupt nesting?
 - Stack usage of library functions should also be considered



Use a Real Time Kernel or Not?

- Ease application development
 - Multi-task applications
 - Provide basic services
 - Provide/handle time critical event efficiently
- Extra Cost
 - More memory (ROM/RAM) space
 - CPU overhead (about 2% - 4%)
 - More money!!
 - Not all RTOS kernels are free!